

GPU FinOps & Cost Optimization - Lab Report

Student: Nguyen Ngoc Tan MSSV: 2A202600190

Notebook: gpu_finops_lab.ipynb

This report summarizes the completed GPU FinOps hands-on lab. The local Docker Compose services emulate a GPU cloud control plane, while the Kaggle notebook calls the gateway through a tunnel to monitor GPU nodes, submit workloads, record billing, evaluate spot savings, run autoscaling decisions, analyze waste, and compare real GPU training options such as FP32 and mixed precision AMP.

Area	Key result / interpretation
Cluster monitoring	The lab monitors GPU nodes, utilization, memory, power, temperature, and idle capacity.
Workload + billing	Workloads are submitted to GPUs and billing events show total cost, savings, and budget status.
Spot management	Spot capacity is cheaper than on-demand and can produce large savings, with preemption risk.
Autoscaling	Autoscaler decisions are based on utilization thresholds and cost-aware policy.
Waste analysis	Idle GPUs create waste cost; the tracker estimates waste percentage and monthly savings.
Recommendations	Main actions are using spot capacity, scheduling non-urgent jobs, and optimizing training.

GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

Gateway Connection Test

```
[OK] Connected to GPU FinOps Lab Gateway
Status code: 200
Response:
{
  "service": "GPU FinOps Lab Gateway",
  "endpoints": {
    "cluster": "/cluster/*",
    "billing": "/billing/*",
    "spot": "/spot/*",
    "autoscaler": "/autoscaler/*",
    "cost": "/cost/*"
  }
}
```

GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

Part 1 - View Cluster Nodes

Status: 200

Preview: {"node-00":[{"node_id":"node-00","gpu_id":0,"gpu_type":"T4","utilization":65.7664290710675,"memory_used_gb":9.083071115282362,"memory_total_gb":16.0,"power_draw_watts":45.51674000268728,"temperature_c

Cluster has 5 nodes

- ```
=====
```
- node-00:  
GPU 0 [T4] Util: 65.8% | Mem: 9.1/16.0GB | Power: 46W | Temp: 81°C  
GPU 1 [T4] Util: 61.5% | Mem: 13.6/16.0GB | Power: 44W | Temp: 66°C
  - node-01:  
GPU 0 [A100] Util: 78.4% | Mem: 40.7/80.0GB | Power: 195W | Temp: 69°C  
GPU 1 [A100] Util: 91.5% | Mem: 61.4/80.0GB | Power: 228W | Temp: 58°C
  - node-02:  
GPU 0 [V100] Util: 78.7% | Mem: 27.3/32.0GB | Power: 226W | Temp: 57°C  
GPU 1 [V100] Util: 80.1% | Mem: 19.4/32.0GB | Power: 237W | Temp: 79°C
  - node-03:  
GPU 0 [T4] Util: 68.8% | Mem: 12.1/16.0GB | Power: 44W | Temp: 77°C  
GPU 1 [T4] Util: 63.6% | Mem: 12.2/16.0GB | Power: 44W | Temp: 59°C
  - node-04:  
GPU 0 [T4] Util: 78.0% | Mem: 13.4/16.0GB | Power: 66W | Temp: 81°C  
GPU 1 [T4] Util: 68.7% | Mem: 9.2/16.0GB | Power: 43W | Temp: 66°C

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 1 - Cluster Metrics Summary

Status: 200

Preview: {"total\_gpus":10,"busy\_gpus":10,"idle\_gpus":0,"avg\_utilization":73.49552444904207,"total\_memory\_used\_gb":218.31,"total\_memory\_capacity\_gb":320.0,"total\_power\_draw\_watts":1172.4,"node\_count":5}

Cluster Metrics

```
=====
Total GPUs: 10
Busy GPUs: 10
Idle GPUs: 0
Avg Utilization: 73.5%
Memory Used: 218.3 GB
Memory Capacity: 320.0 GB
Total Power Draw: 1172 W
Node Count: 5
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

## Part 2 - Workload Submission

Submitting workloads...

```
train-resnet-001: queued -> queued
train-bert-002: queued -> queued
inference-api-003: queued -> queued
train-llm-004: queued -> queued
```

Updated metrics:

Busy GPUs: 10/10 | Utilization: 73.5%

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

## Part 2 - Billing Summary

Recording billing events...

train-resnet-001 [ON-DEMAND]: \$0.0292 (saved \$0.0000)

train-bert-002 [ON-DEMAND]: \$0.6117 (saved \$0.0000)

inference-api-003 [SPOT]: \$0.0035 (saved \$0.0082)

train-llm-004 [SPOT]: \$0.5505 (saved \$1.2845)

Billing Summary:

Total Cost: \$2.5589

Total Savings: \$2.7078

Budget Used: 2.6%

Alert Status: OK

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 3 - Spot Pricing

Status: 200

Preview: {"T4":{"on\_demand\_price":0.35,"current\_spot\_price":0.2644,"discount\_pct":24.5,"availability":"medium"},"A100":{"on\_demand\_price":3.67,"current\_spot\_price":2.58,"discount\_pct":29.7,"availability":"medium"}}

Current Spot Pricing

| GPU Type | On-Demand | Spot Price | Discount | Availability |
|----------|-----------|------------|----------|--------------|
| T4       | \$0.35    | \$0.2644   | 24.5     | % medium     |
| A100     | \$3.67    | \$2.5800   | 29.7     | % medium     |
| V100     | \$2.48    | \$1.7843   | 28.1     | % medium     |

## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 3 - Spot Request

Requesting Spot Instances...

[OK] spot-t4-001 (T4): granted

[OK] spot-t4-002 (T4): granted

[OK] spot-a100-001 (A100): granted



## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 3 - Spot Preemption and Savings

✂ Simulating spot preemption event...

Preempted instances: 1

Still active: 5

[!] Preempted:

- opt-spot-2 (ran for 4431s, 60s warning)

Spot Savings Report:

Spot cost: \$0.2590

On-demand equiv: \$0.8634

Total saved: \$0.6043 (70.0%)

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 4 - Autoscaler Policy

```
Current Autoscaling Policy:
 scale_up_threshold: 70.0
 scale_down_threshold: 25.0
 cooldown_seconds: 30
 max_nodes: 10
 min_nodes: 2
 preferred_gpu_type: T4
 cost_aware: True
```

```
Updating policy...
[OK] Policy updated
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 4 - Autoscaler Evaluation

Evaluating autoscaling decision...

† Action: SCALE\_UP

Reason: Utilization 73.5% > threshold 70.0%

Current utilization: 73.5%

Nodes: 5 -> 6

Running 5 evaluation cycles...

Cycle 1: no\_action | Util: 61.2% | Nodes: 6->6

Cycle 2: no\_action | Util: 61.2% | Nodes: 6->6

Cycle 3: no\_action | Util: 61.2% | Nodes: 6->6

Cycle 4: no\_action | Util: 61.2% | Nodes: 6->6

Cycle 5: no\_action | Util: 61.2% | Nodes: 6->6

## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 5 - Cost Snapshots

Taking cost snapshots...

|             |                  |  |                 |  |            |
|-------------|------------------|--|-----------------|--|------------|
| Snapshot 1: | Total=\$0.041944 |  | Idle=\$0.001944 |  | Waste=4.6% |
| Snapshot 2: | Total=\$0.041944 |  | Idle=\$0.001944 |  | Waste=4.6% |
| Snapshot 3: | Total=\$0.041944 |  | Idle=\$0.001944 |  | Waste=4.6% |
| Snapshot 4: | Total=\$0.041944 |  | Idle=\$0.001944 |  | Waste=4.6% |
| Snapshot 5: | Total=\$0.041944 |  | Idle=\$0.001944 |  | Waste=4.6% |

## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

# Part 5 - Waste Report

### WASTE ANALYSIS REPORT

=====

|                         |            |
|-------------------------|------------|
| Average Waste:          | 12.1%      |
| Total Idle Cost:        | \$0.046996 |
| Total Cost:             | \$0.401944 |
| Potential Monthly Save: | \$1218.14  |
| Severity:               | LOW        |

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

# Part 5 - Recommendations

### COST OPTIMIZATION RECOMMENDATIONS

- ```
=====
```
1. [MEDIUM] USE_SPOT
Switch fault-tolerant workloads to spot instances for 60-70% savings.
Estimated savings: 65.0%
 2. [LOW] SCHEDULING
Schedule non-urgent training jobs during off-peak hours for lower spot prices.
Estimated savings: 20.0%

GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

Part 5 - Full Dashboard

GPU FinOps DASHBOARD


- ```
=====
```
- CLUSTER: 12 GPUs across 6 nodes  
Utilization: 61.2% | Busy: 10 | Idle: 2
  - BILLING: \$2.5589 / \$100.00 budget  
Alert: OK | Savings: \$2.7078
  - SPOT: Saved \$0.6294 (70.0%)
  - WASTE: 12.1% | Severity: LOW

## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

# Part 6 - Cost Breakdown Visualization

Chart saved as finops\_cost\_breakdown.png



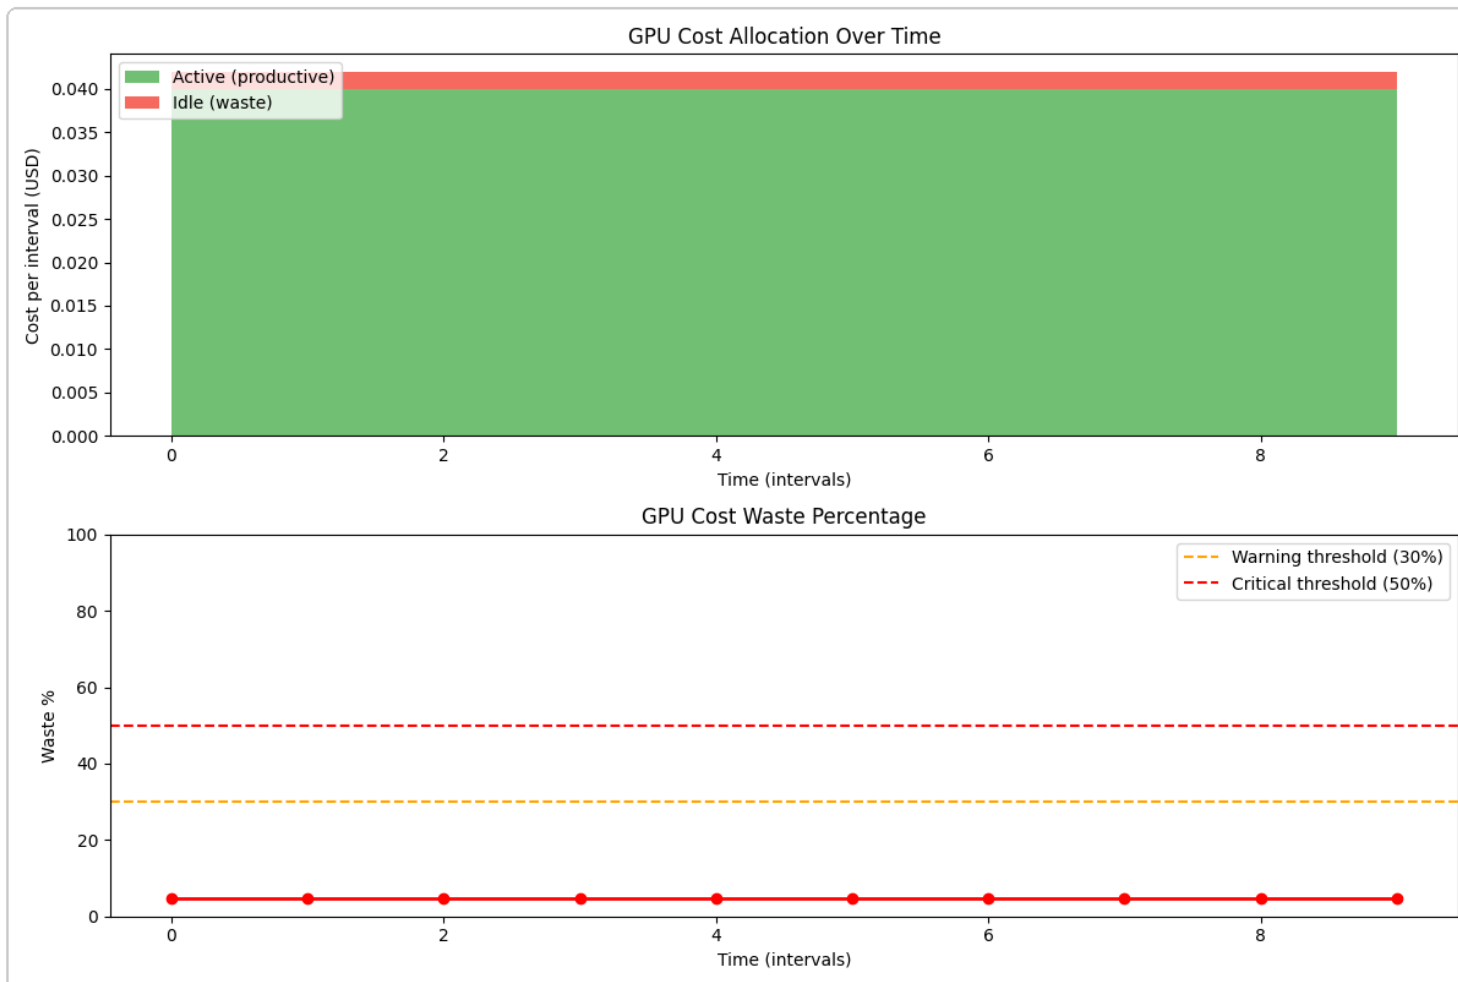


## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 6 - Time-series Cost Tracking

Collecting time-series data (10 snapshots)...



## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

# Part 7 - Full FinOps Optimization Workflow

### FULL FINOPS OPTIMIZATION WORKFLOW

```
=====

1[] Initial cluster state:
 GPUs: 12 | Util: 61.2% | Idle: 2

2[] Submitting heavy workloads...
 After load: Util: 74.8% | Busy: 12/12

3[] Autoscaler evaluation:
 Decision: scale_up - Utilization 74.8% > threshold 70.0%

4[] Cost analysis:
 Total cost/interval: $0.043889
 Waste: 4.4%

5[] Recommendations:
 [MEDIUM] USE_SPOT: savings ~65.0%
 [LOW] SCHEDULING: savings ~20.0%

6[] Applying optimization: Switch to spot instances...
 Spot savings: $0.0451 (70.0%)

7[] Final billing:
 Total spend: $2.7280
 Total saved: $2.8302
 Budget: 2.7% used

[OK] Workflow complete!
```

## GPU FinOps Lab - Student Information

Họ và tên: Nguyen Ngoc Tan | MSSV: 2A202600190

## Part 8 - Real GPU Detection

### Real GPU Detected

Name: Tesla T4  
Memory: 15.6 GB  
Type: T4  
Pricing: \$0.35/hr (on-demand)  
CUDA: 12.8  
pynvml: available

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8 - GPU Metrics Diagnostic

```
=====
GPU METRICS DIAGNOSTIC
=====
```

```
1. pynvml available: True
 pynvml works! GPU util=0%, mem=472/16106 MB
 Power: 10.7W
 Temp: 42C
```

```
2. get_gpu_metrics() test:
 gpu_util_pct: 0.0
 mem_total_mb: 16106.12736
 mem_used_mb: 472.055808
 mem_util_pct: 0.0
 power_watts: 10.737
 temp_c: 42.0
```

```
Method: pynvml
```

```
Ready for training.
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8 - FP32 Training Summary

```
=====
EXPERIMENT 1: FP32 Training (Baseline)
```

```
=====
```

```
Epoch 1/3 | Loss: 2.0593 | Acc: 26.8% | Time: 40.2s | Samples: 40
Epoch 2/3 | Loss: 1.4131 | Acc: 48.1% | Time: 40.9s | Samples: 40
Epoch 3/3 | Loss: 1.1214 | Acc: 59.8% | Time: 45.9s | Samples: 40
```

```
Total samples collected: 120
```

```
FP32 Summary:
```

```
Total time: 127.1s
Peak memory: 0.82 GB
Avg GPU util: 96.0%
Avg power: 67.5W
Avg temperature: 68.5C
Max GPU util: 99.0%
Estimated cost: $0.012358
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

## Part 8 - AMP Training Summary

```
=====
EXPERIMENT 2: Mixed Precision (AMP) Training
=====
```

```
Epoch 1/3 | Loss: 1.8733 | Acc: 32.3% | Time: 20.0s | Samples: 40
Epoch 2/3 | Loss: 1.3722 | Acc: 49.7% | Time: 19.4s | Samples: 40
Epoch 3/3 | Loss: 1.0850 | Acc: 61.1% | Time: 19.7s | Samples: 40
```

Total samples collected: 120

### AMP Summary:

```
Total time: 59.1s
Peak memory: 0.60 GB
Avg GPU util: 91.4%
Avg power: 65.5W
Avg temperature: 76.2C
Max GPU util: 94.0%
Estimated cost: $0.005747
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8 - FP32 vs AMP Comparison

=====

FINOPS COMPARISON: FP32 vs Mixed Precision (AMP)

=====

| Metric           | FP32       | AMP        | Improvement      |
|------------------|------------|------------|------------------|
| Total Time       | 127.1      | 59.1       | 2.15x faster     |
| Peak Memory (GB) | 0.82       | 0.60       | 0.22 GB saved    |
| Cost (USD)       | \$0.012358 | \$0.005747 | \$0.006611 saved |
| Cost Saving %    | ---        | ---        | 53.5%            |
| Avg GPU Util %   | 96.0       | 91.4       |                  |
| Avg Power (W)    | 67.5       | 65.5       |                  |

--- Extrapolated Savings at Scale ---

1 day training: FP32=\$8.40 vs AMP=\$3.91 -> SAVE \$4.49

1 week training: FP32=\$58.80 vs AMP=\$27.34 -> SAVE \$31.46

1 month training: FP32=\$252.00 vs AMP=\$117.19 -> SAVE \$134.81



## Part 8 - Real GPU Cost Report

```
=====
REPORTING REAL GPU COSTS TO FINOPS GATEWAY
=====

GET https://cfeac3468cad39.lhr.life/
Status: 200
Content-Type: application/json
Response preview: {"service": "GPU FinOps Lab Gateway", "endpoints": {"cluster": "/cluster/*", "billing": "/bil
ling/*", "spot": "/spot/*", "autoscaler": "/autoscaler/*", "cost": "/cost/*"}}

Gateway health: {'service': 'GPU FinOps Lab Gateway', 'endpoints': {'cluster': '/cluster/*', 'billing':
'/billing/*', 'spot': '/spot/*', 'autoscaler': '/autoscaler/*', 'cost': '/cost/*'}}

POST https://cfeac3468cad39.lhr.life/billing/record
Status: 200
Content-Type: application/json
Response preview: {"workload_id": "real-gpu-resnet18-
fp32", "gpu_type": "T4", "gpu_count": 1, "duration_seconds": 127.10727620124817, "is_spot": false, "rate_per_hour"
: 0.35, "total_cost_usd": 0.0124, "savings_usd": 0, "project": "real-gpu-lab", "timestamp": 1778650391.4616735}

FP32 workload reported:
Cost: $0.012400 | Rate: $0.3500/hr

POST https://cfeac3468cad39.lhr.life/billing/record
Status: 200
Content-Type: application/json
Response preview: {"workload_id": "real-gpu-resnet18-
amp", "gpu_type": "T4", "gpu_count": 1, "duration_seconds": 59.10841774940491, "is_spot": true, "rate_per_hour": 0.
105, "total_cost_usd": 0.0017, "savings_usd": 0.004, "project": "real-gpu-lab", "timestamp": 1778650391.952357}

AMP workload reported (as spot):
Cost: $0.001700 | Saved: $0.004000

POST https://cfeac3468cad39.lhr.life/cluster/workloads/submit
Status: 200
Content-Type: application/json
Response preview: {"status": "running", "assigned": [{"node_id": "node-05", "gpu_id": 0}], "workload_id": "real-
gpu-resnet18-fp32"}

POST https://cfeac3468cad39.lhr.life/cluster/workloads/submit
Status: 200
Content-Type: application/json
Response preview: {"status": "running", "assigned": [{"node_id": "node-05", "gpu_id": 1}], "workload_id": "real-
gpu-resnet18-amp"}

--- Updated FinOps Billing (incl. real GPU) ---

GET https://cfeac3468cad39.lhr.life/billing/summary
Status: 200
Content-Type: application/json
Response preview: {"project": "real-gpu-lab", "total_cost_usd": 0.0141, "total_savings_usd": 0.004, "budget_usd
": 100.0, "budget_remaining_usd": 99.9859, "budget_utilization_pct": 0.01, "total_workloads": 2, "cost_by_gpu_typ
e": {"T4": {"cost": 0.0141, "hours": 0.05172658165295919, "workloads": 2}}, "alert": "OK"}
Project: real-gpu-lab
Total Cost: $0.014100
Total Savings: $0.004000
Workloads: 2

POST https://cfeac3468cad39.lhr.life/cost/snapshot
Status: 200
Content-Type: application/json
Response preview: {"timestamp": 1778650394.1873095, "node_costs": {"node-
00": {"total_cost": 0.001944, "idle_cost": 0, "gpu_count": 2}, "node-
01": {"total_cost": 0.020389, "idle_cost": 0, "gpu_count": 2}, "node-
02": {"total_cost": 0.013778, "idle_cost": 0.006889, "gpu_count": 2}, "node-
03": {"total_cost": 0.001944, "idle_cost": 0, "gpu_count": 2}, "node-
04": {"total_cost": 0.001944, "idle_cost": 0, "gpu_count": 2}, "node-
05": {"total_cost": 0.001944, "idle_cost": 0, "gpu_count": 2}, "node-
06": {"total_cost": 0.001944, "idle_cost": 0.001944, "gpu_count": 2}}, "tota

Cost snapshot taken: waste=20.1%

--- FINAL DASHBOARD (Mock + Real GPU) ---

GET https://cfeac3468cad39.lhr.life/cost/dashboard
Status: 200
Content-Type: application/json
Response preview: {"cluster_metrics": {"total_gpus": 14, "busy_gpus": 11, "idle_gpus": 3, "avg_utilization": 59.5
9311393068521, "total_memory_used_gb": 223.54, "total_memory_capacity_gb": 384.0, "total_power_draw_watts": 112
9.46, "node_count": 7}, "billing_summary": {"project": "default", "total_cost_usd": 2.728, "total_savings_usd": 2.
8302, "budget_usd": 100.0, "budget_remaining_usd": 97.272, "budget_utilization_pct": 2.73, "total_workloads": 18,
"cost_by_gpu_type": {"T4": {"cost": 0.40360000000000007, "hours": 1.9000000000000004, "workloads": 14}, "A
Total Platform Cost: $2.7280
Total Savings: $2.8302
Budget Utilization: 2.7%
Alert: OK
```





## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8 - Real GPU Visualizations

Real GPU Telemetry During Training

Charts saved: cost\_per\_epoch.png, real\_gpu\_telemetry.png



## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8.5 - Multi-GPU Cost Analysis

```
=====
EXERCISE 8.5.1: Multi-GPU Cost Analysis
=====
```

|   | gpu_count | gpu_type | speedup | efficiency_pct | time_hours | total_cost_usd | \ |
|---|-----------|----------|---------|----------------|------------|----------------|---|
| 0 | 1         | A100     | 1.00    | 100.0          | 2.000000   | 7.340000       |   |
| 1 | 2         | A100     | 1.84    | 92.0           | 1.086957   | 7.978261       |   |
| 2 | 4         | A100     | 3.36    | 84.0           | 0.595238   | 8.738095       |   |
| 3 | 8         | A100     | 6.08    | 76.0           | 0.328947   | 9.657895       |   |

|   | cost_per_speedup |
|---|------------------|
| 0 | 7.340000         |
| 1 | 4.336011         |
| 2 | 2.600624         |
| 3 | 1.588470         |

[OK] Best cost option: 1 GPU(s)  
Time: 2.00h | Cost: \$7.34 | Efficiency: 100.0%

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8.5 - Project Cost Forecast

```
=====
EXERCISE 8.5.2: Project Cost Forecasting
=====
```

|   | phase                 | gpu_type | gpu_count | duration_hours | base_cost_usd | \ |
|---|-----------------------|----------|-----------|----------------|---------------|---|
| 0 | Data Preparation      | T4       | 1         | 40             | 14.0          |   |
| 1 | Model Training        | A100     | 4         | 120            | 1761.6        |   |
| 2 | Hyperparameter Tuning | A100     | 8         | 60             | 1761.6        |   |
| 3 | Model Evaluation      | T4       | 2         | 20             | 14.0          |   |

|   | low_estimate_usd | high_estimate_usd | uncertainty_pct |
|---|------------------|-------------------|-----------------|
| 0 | 11.90            | 16.10             | 15.0            |
| 1 | 1321.20          | 2202.00           | 25.0            |
| 2 | 1233.12          | 2290.08           | 30.0            |
| 3 | 12.60            | 15.40             | 10.0            |

```
[OK] Base cost: $3,551.20
[OK] Contingency: $710.24
[OK] Expected total: $4,261.44
[OK] 95% confidence: $3,573.51 - $4,949.37
```

## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8.5 - Optimization Opportunity Analysis

```
=====
EXERCISE 8.5.3: Advanced Optimization Opportunity Analysis
=====
Baseline cost: $1,468.00

 name savings_pct estimated_savings_usd \
0 Switch to Mixed Precision (AMP) 25.0 367.0
1 Optimize Batch Size 15.0 220.2
2 Implement Early Stopping 20.0 293.6
3 Use Spot Instances 60.0 880.8
4 Switch to More Efficient GPU Type 40.0 587.2

implementation_effort risk_level priority_score dependencies
0 LOW LOW 367.000000 None
1 LOW LOW 220.200000 None
2 MEDIUM LOW 146.800000 None
3 MEDIUM HIGH 146.800000 None
4 HIGH MEDIUM 97.866667 None

Roadmap:

 step strategy incremental_savings_usd \
0 1 Switch to Mixed Precision (AMP) 367.0000
1 2 Optimize Batch Size 165.1500
2 3 Implement Early Stopping 187.1700
3 4 Use Spot Instances 449.2080
4 5 Switch to More Efficient GPU Type 119.7888

 remaining_cost_usd
0 1101.0000
1 935.8500
2 748.6800
3 299.4720
4 179.6832

[OK] Optimized cost: $179.68
[OK] Total savings: $1,288.32 (87.8%)
```


## GPU FinOps Lab - Student Information

Ho va ten: Nguyen Ngoc Tan | MSSV: 2A202600190

### Part 8.5 - Integrated Cost Dashboard

#### EXERCISE 8.5.4: Integrated Cost Dashboard

[OK] Saved dashboard as advanced\_finops\_dashboard.png



Part 8.5 - Challenge Strategy

```
=====
SOLUTION: Cost Optimization Strategy
=====

1. BASELINE COST
GPU: 8x A100
Duration: 200 hours
Rate: $3.67/GPU-hour
Baseline cost: $5,872.00

2. MULTI-GPU COST ANALYSIS
 4 GPUs | eff=0.93 | time=432.4h | cost=$6,348.11 | deadline=MISS
 6 GPUs | eff=0.88 | time=304.8h | cost=$6,710.86 | deadline=OK
 8 GPUs | eff=0.82 | time=242.4h | cost=$7,117.58 | deadline=OK
10 GPUs | eff=0.78 | time=206.5h | cost=$7,576.77 | deadline=OK
12 GPUs | eff=0.72 | time=183.9h | cost=$8,099.31 | deadline=OK
16 GPUs | eff=0.62 | time=160.0h | cost=$9,395.20 | deadline=OK

Recommended GPU count: 6 GPUs ($6,710.86, 304.8h)

3. SELECTED OPTIMIZATION STRATEGIES
- Mixed precision / AMP: 35.0% savings, risk=LOW
 Reason: Reduces training time and memory use while preserving target accuracy.
- Spot instances with checkpointing: 55.0% savings, risk=MEDIUM
 Reason: Uses cheaper capacity while checkpointing protects against preemption.
- Deadline-aware GPU sizing: 0.0% savings, risk=LOW
 Reason: Chooses the cheapest GPU count that still meets the 2-week deadline.
- Idle shutdown and job queue scheduling: 10.0% savings, risk=LOW
 Reason: Reduces idle GPU waste between experiments.
- Frequent checkpointing and resume: 5.0% savings, risk=LOW
 Reason: Reduces lost work during spot preemption.

4. CUMULATIVE SAVINGS
Baseline cost: $5,872.00
Deadline-aware plan cost: $6,710.86
Optimized cost: $1,678.30
Savings vs baseline: $4,193.70 (71.4%)

5. COST FORECAST WITH UNCERTAINTY
Expected cost: $1,678.30
Low estimate: $1,342.64
High estimate: $2,013.96
Confidence interval: ±20%

6. CONSTRAINT CHECK
Under budget: True ($1,678.30 <= $5,000.00)
Meets deadline: True (304.8h <= 336h)
Preemption risk OK: True
Accuracy target OK: True

7. STRATEGY SUMMARY
Decision: APPROVED
Recommended plan:
- Use 6x A100
- Switch from FP32 to AMP/mixed precision
- Use spot instances with checkpointing
- Enable idle shutdown and queue-based scheduling
- Track cost daily through FinOps dashboard

Justification:
The original plan costs $5,872.00, which is above the $5,000.00 budget. The deadline-aware GPU plan uses 6x A100 and finishes in 304.8h, within the 336h deadline. After AMP, spot, idle shutdown, and checkpointing, the expected cost is $1,678.30, while keeping preemption risk within the MEDIUM limit and preserving the minimum accuracy target of 0.95.
```

## Conclusion

The lab demonstrates the core GPU FinOps loop: gain visibility into cluster utilization, associate workloads with cost, choose cheaper capacity such as spot where safe, use autoscaling to control idle capacity, and use waste analysis to prioritize optimization work. The real GPU training section also shows that code-level optimizations such as mixed precision can reduce runtime and cost, complementing infrastructure-level optimizations.

Deliverables generated in this package: `report.pdf`, `screenshots/`, `generated_charts/`, and `notebook/gpu_finops_lab.ipynb`.